

# DSO 545: Statistical Computing and Data Visualization

*Data Manipulation with dplyr : Analyzing Flights Data (hflights)*

*Fall 2018*

## LAB 03

### The Five Verbs of dplyr

1. The **dplyr** package contains five key data manipulation functions, also called verbs:
  - **select()**, which returns a subset of the columns,
  - **filter()**, that is able to return a subset of the rows,
  - **arrange()**, that reorders the rows according to single or multiple variables,
  - **mutate()**, used to add columns from existing data,
  - **summarise()**, which reduces each group to a single row by calculating aggregate measures.

What order of operations should we use to find the average value of the **ArrDelay** (arrival delay) variable for all American Airline flights in the **hflights** tbl?

### Manipulating Variables (Select and Mutate)

2. Return a copy of **hflights** that contains the four columns related to delay (**ActualElapsedTime**, **AirTime**, **ArrDelay**, **DepDelay**).
3. Return a copy of **hflights** containing the columns **Origin** up to **Cancelled**.
4. Find the most concise way to select: columns **Year** up to and including **DayOfWeek**, columns **ArrDelay** up to and including **Diverted**.
5. **dplyr** comes with a set of helper functions that can help you select variables. These functions find groups of variables to select, based on their names. **dplyr** provides 6 helper functions, each of which only works when used inside **select()**.
  - **starts\_with("X")**: every name that starts with "X",
  - **ends\_with("X")**: every name that ends with "X",
  - **contains("X")**: every name that contains "X",
  - **matches("X")**: every name that matches "X", which can be a regular expression,
  - **num\_range("x", 1:5)**: the variables named x01, x02, x03, x04 and x05,
  - **one\_of(x)**: every name that appears in x, which should be a character vector.

Use **select** and a helper function to return a tbl copy of **hflights** that contains just **ArrDelay** and **DepDelay**.

6. Use a combination of helper functions and variable names to return the **UniqueCarrier**, **FlightNum**, **TailNum**, **Cancelled**, and **CancellationCode** columns of **hflights**.
7. Which variables in **hflight** do you think count as a plane's "ground time"? Use **mutate()** to add these variables together and save them as **GroundTime**.

### Manipulating Observations (Filter and Arrange)

When manipulating observations, we should know the following operations:

- **x < y**, TRUE if x is less than y
- **x <= y**, TRUE if x is less than or equal to y

- `x == y`, TRUE if x equals y
  - `x != y`, TRUE if x does not equal y
  - `x >= y`, TRUE if x is greater than or equal to y
  - `x > y`, TRUE if x is greater than y
  - `x %in% c(a, b, c)`, TRUE if x is in the vector `c(a, b, c)`
8. Return a copy of all flights that traveled 3000 miles or more.
  9. Return a copy of all flights flown by one of American(AA), Alaska (AS), or JetBlue (B6) airlines.
  10. Return a copy of all flights where taxi-ing took longer than flying.
  11. Return a copy of all cancelled weekend flights
  12. Arrange according to carrier and decreasing departure delays.
  13. Arrange flights by total delay (normal order).
  14. Filter out flights leaving to DFW before 8am and arrange according to decreasing AirTime

## Manipulating Groups of Observation (`summarize` and `group_by`)

15. Determine the shortest and longest distance flown and save statistics to `min_dist` and `max_dist` resp.
16. Determine the longest distance for diverted flights, save statistic to `max_div`.
17. `dplyr` provides several helpful aggregate functions of its own, in addition to the ones that are already defined in R. These include:
  - `first(x)` - The first element of vector x.
  - `last(x)` - The last element of vector x.
  - `nth(x, n)` - The nth element of vector x.
  - `n()` - The number of rows in the data.frame or group of observations that `summarise()` describes.
  - `n_distinct(x)` - The number of unique values in vector x.

Create a table with the following variables (and variable names): the total number of observations in `hflights` (`n_obs`), the total number of carriers that appear in `hflights` (`n_carrier`), the total number of destinations that appear in `hflights` (`n_dest`), and the destination of the flight that appears in the 100th row of `hflights` (`dest100`).

18. Use Piping: (1) Take the `hflights` data set and then, (2) Add a variable named `diff` that is the result of subtracting `TaxiIn` from `TaxiOut`, and then (3) pick all of the rows whose `diff` value does not equal NA, and then (4) summarise the data set with a value named `avg` that is the mean `diff` value.
19. Use Piping: Define a data set named `d` that contains just the `Dest`, `UniqueCarrier`, `Distance`, and `ActualElapsedTime` columns of `hflights` as well an additional variable: `RealTime` which is equal the actual elapsed time plus 100 minute.

## Grouped Summaries using '`group_by()`' and '`summarize()`'

20. For each destination, find the number of flights, the mean distance travelled, and the mean arrival delay.
21. Investigate visually if there is any relationship between distance travelled and arrival delays. You can get rid of the outliers if any.
22. Use `group_by()` and `summarise()` to compare the individual carriers. For each carrier, count the total number of flights flown by the carrier (`n_flights`), the total number of cancelled flights (`n_canc`), and the average arrival delay of the flights whose delay does not equal NA (`avg_delay`). Once you've calculated these results, `arrange()` the carriers from low to high by their average arrival delay. Use number of flights cancelled to break any ties. Which airline scores best based on these statistics?